A person wearing a dark hoodie and a cap, sitting at a desk and working on a laptop.

BENCHMARKING and PROFILING in RUBY

BENCHMARKING and PROFILING in RUBY



SUMMER WORKSHOP TOUR

buy @ speedshop.co

- Minneapolis/RailsConf Friday May 3rd.
Hyatt Regency (same as the conference hotel).
- New York City Monday June 3rd. Chelsea.
- New York City Wednesday June 5th.
Chelsea.
- New York City Friday June 7th. Special Topic: ActiveRecord Deep Dive.
- Atlanta, GA Monday June 17th.
- Austin, TX Wednesday June 19th.
- Dallas, TX Friday June 21st.
- Denver, CO Monday June 24th.
- Chicago, IL Wednesday June 26th.
- Boston, MA Monday July 15th.
- Washington DC Wednesday July 17th.
- Philadelphia, PA Friday July 19th.
- Portland, OR Monday July 29th.
- Seattle, WA Wednesday July 31st.
- San Francisco, CA Friday August 2nd.
- San Francisco, CA Monday August 5th.
- San Francisco, CA Wednesday August 7th. Special Topic: ActiveRecord

FLIGHT PLAN

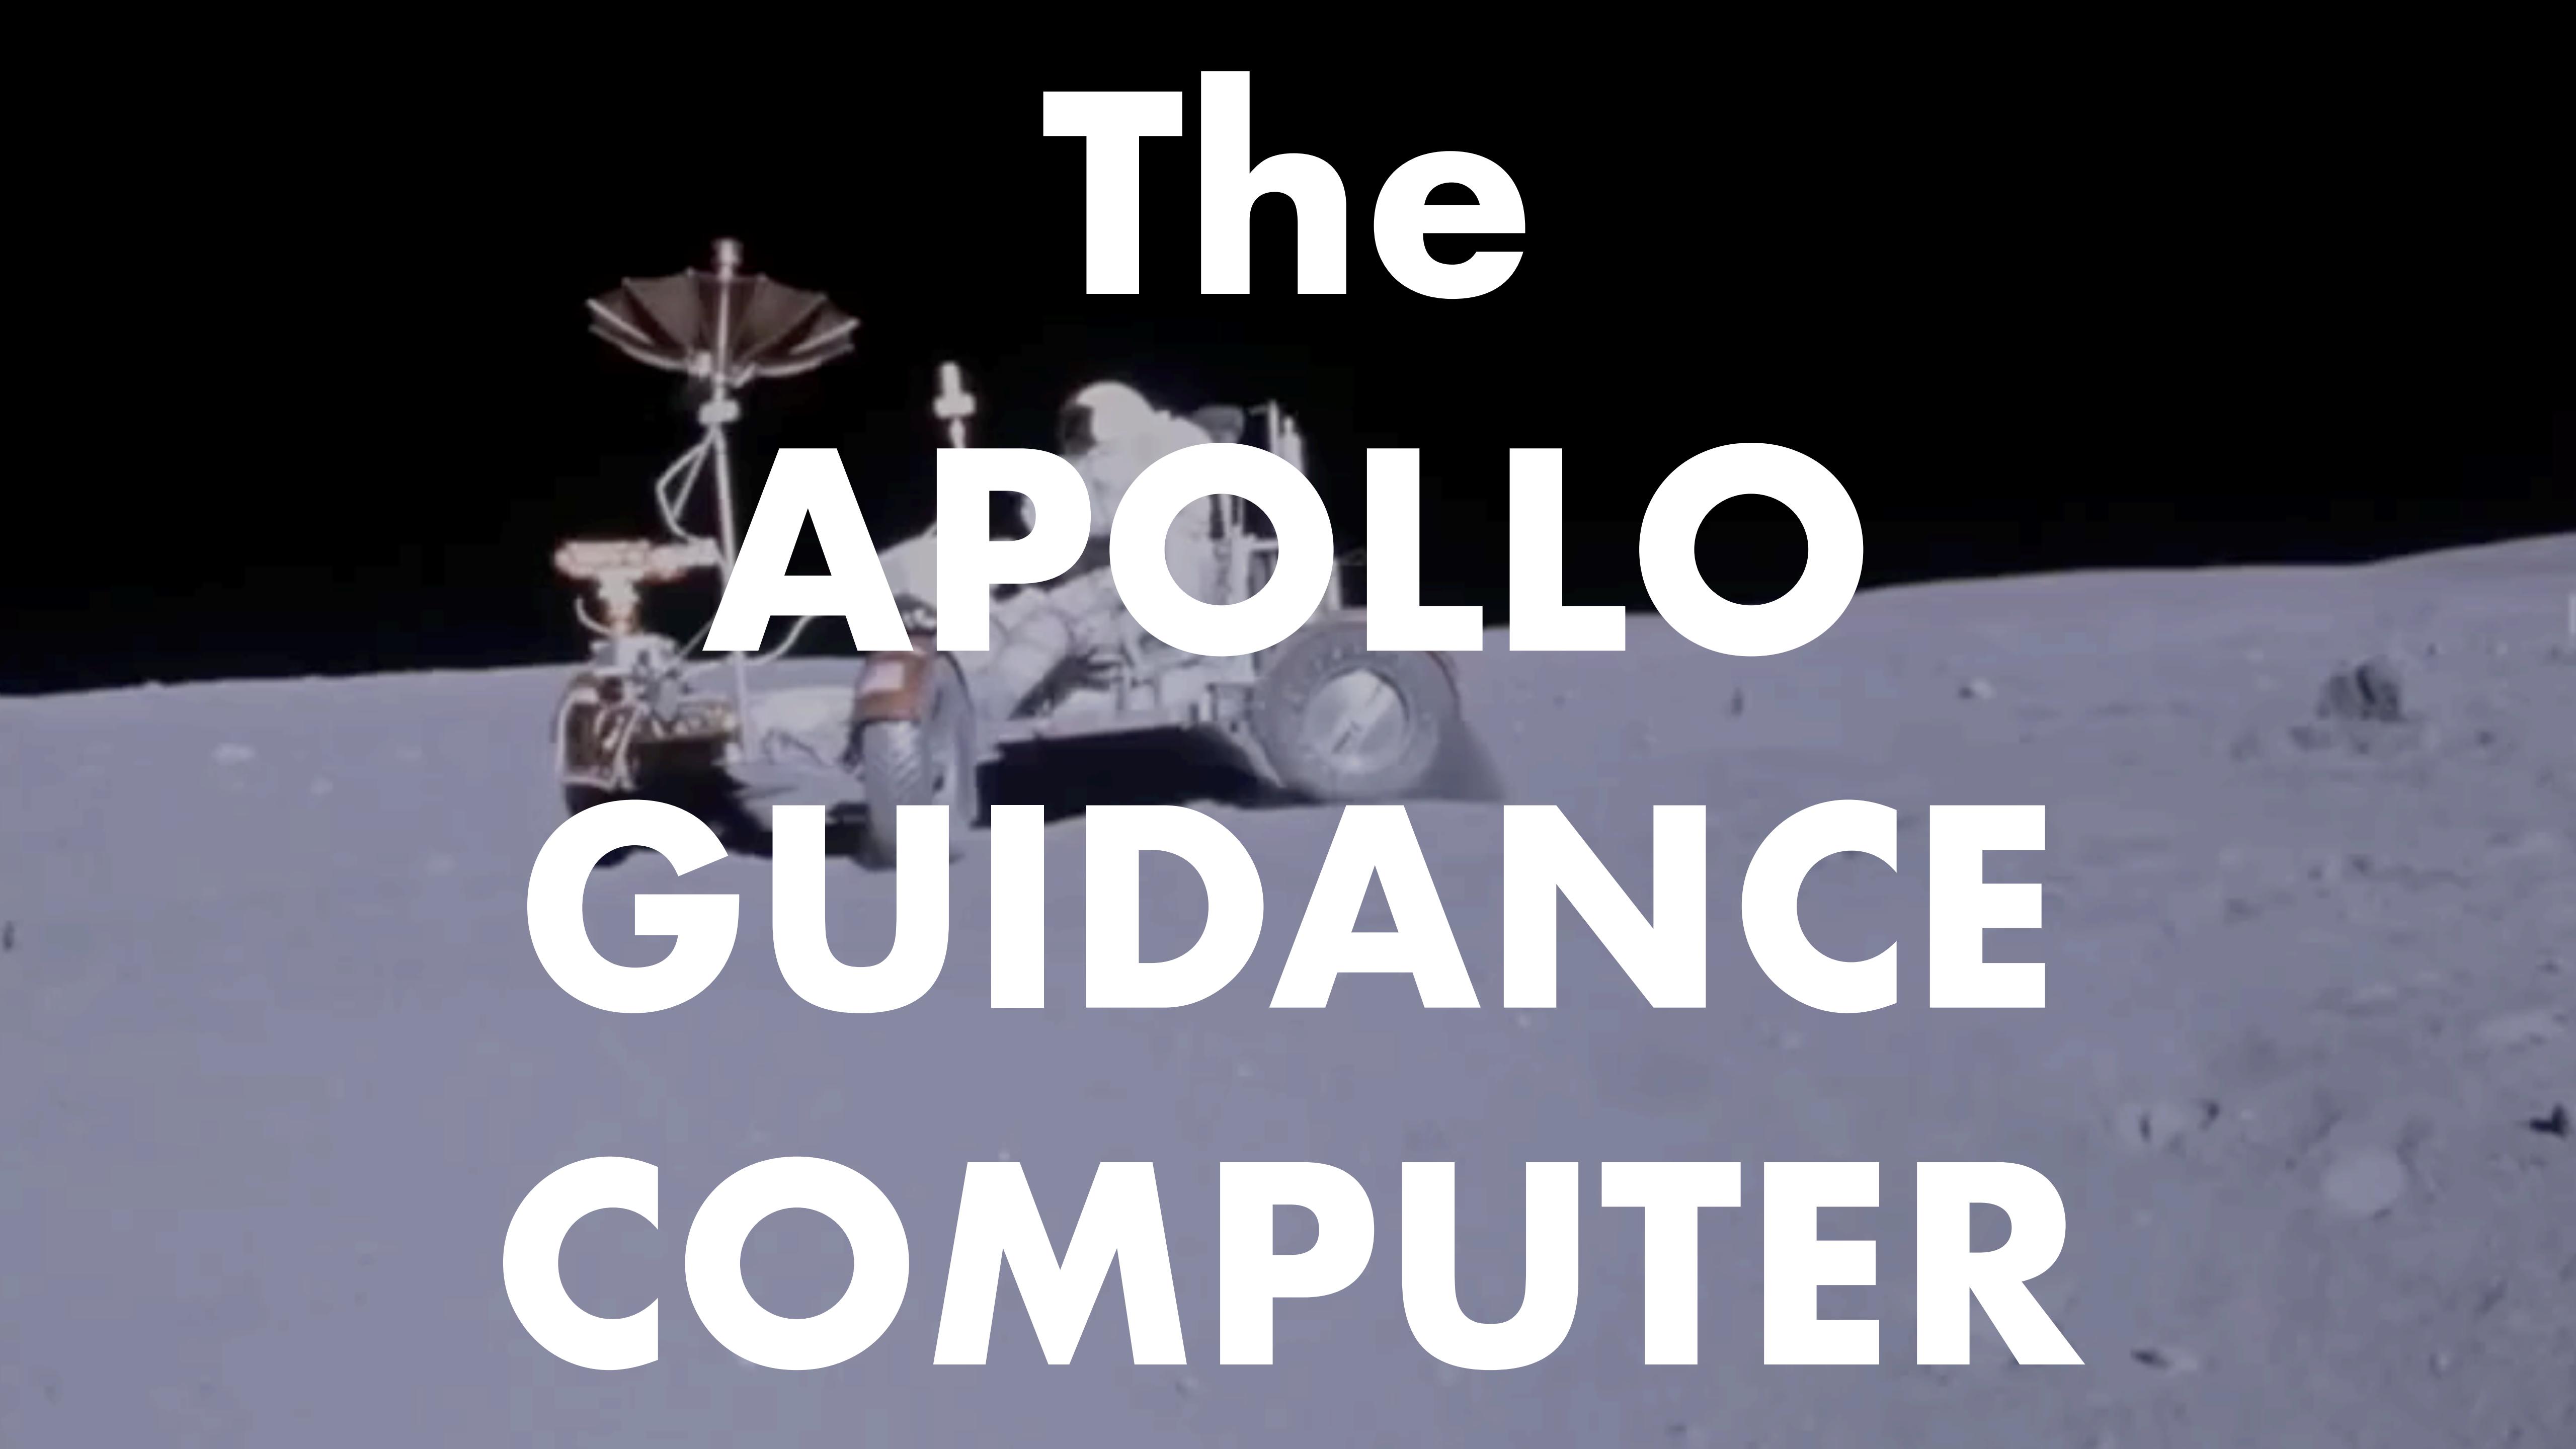
Why Measure?

Benchmarking vs. Profiling

Profiling Jumpstart

Benchmarking Jumpstart

The APOLLO GUIDANCE COMPUTER





FIRST MAN (2018)

04 06 38 30 CDR (EAGLE)

It's a 1202.

04 06 38 32 LMP (EAGLE)

1202.

04 06 38 48 CDR (EAGLE)

Give us a reading on the
1202 PROGRAM ALARM.

04 06 38 53 CC

Roger. We got - We're GO
on that alarm



APPLICABLE TO: IN DESCENT, AVERAGE-G ON

ALARM CODE	TYPE	PRE-HANDOVER CAPABILITY	HANDOVER CAPABILITY
0105 MK ROUT. BUSY 00430 CANT INTG. SV.	P00D00	*	PGNCS GUIDANCE NO/GO
01103 CCSHOLE-PROG.BUG 01204 NEG. WAITLIST 01206 DSKY, TWO USERS 01302 NEG. SQ. ROOT 01501 DSKY, PROG. BAD 01502 DSKY, PROG. BUG 00607 LAHB, NO SOLN	"	* PGNCS GUID. LOST, * PGNCS/AGS ABRT/ABRT STG * (Decision how on current rules) * (NO LR DATA)	(PGNCS GO FOR TAPE METERS, CROSS-POINTERS, CONTROL, ABORTING) (NO LR DATA)
"O.F."=Overflow,to many- CONTINUING ← OCCURRENCE OF?		DUTY CYCLE MAY DEGRADE PGNCS (AGS CONTROL MAY HELP-SF E BELOW) [WATCH FOR OTHER CUES]	SAME AS LEFT
01109 DELAY ROUT. O.V. 01201 EXECT. O.F.(VAC) 01202 EXECT. O.F.(JOBS) 01203 EXECT. O.F.(TASKS) 01207 EXECT. O.F.(MRVS) 01210 TWO USERS 01211 MRK ROUT. INTKPT 02000 DAP O.F.	BALLOUT	PGNCS CONDITION UNKNOWN DSKY MAY BE LOCKED UP, DUTY CYCLE MAY BE UP TO POINT OF MISSING SOME FUNCTIONS (NAV, LAST TO DIE) SWITCH TO AGS (FOLLOW ERR NEEDLES) MAY HELP (REDUCES PGNCS DUTY CYCLE SIGNIF.)	(except "other cues" which would otherwise be cause for ABORT PROBABLY ARENT, INSTEAD IT WOULD BE PGNCS GUIDANCE NO/GO - COMPLETELY MANUAL LANDING IN AGS.)
<u>ISS WARNINGS WITH:</u> 00177 PIPA FAIL 03777 CDU FAIL 09777 PIPA,CDU FAIL 07777 IMU FAIL 10777 PIPA,IMU FAIL 13777 CDU,IMU FAIL 14777 PIPA,CDU,IMU FAIL	LIGHT ONLY	PIPA/CDU/IMU FAIL DISCRETES PRESENT (Other mission rules suffice; alarm may help point to what rule will be broken)	Same as left
00214 IMU TURNOFF	LIGHT ONLY	* AGS ABRT/ABRT STAGE * AGS ABRT/ABRT STAGE *	SWITCH TO AGS PGNCS NO/GO on GND/C (POSS. NO/GO on NAV.)
01107 E-Mem. Destroyed	FRESH SRT	* AGS ABRT/ABRT STAGE * AGS ABRT/ABRT STAGE *	SWITCH TO AGS PGNCS NO/GO! (IMU as ref. okay)
<u>CONTINUING</u> ← 00402 BADGUID. CMDS	LIGHT ONLY	* IF ALARM DOESN'T STOP: * SAME AS "P00D00's" (ABRT!)	IF ALARM DOESN'T STOP: Same as "P00D00"
<u>CONTINUING</u> ← 01406 GUID. NO SOLN 01410 GUID O.V.	LIGHT ONLY	PGNCS GUID. NO/GO AS LONG AS ALARM OCCURRING (ATT. HOLD, CONST. GTC, CONT. OK) (ABRT WILL PROB. COME FROM CURRENT RULES e.g. GTC vs. V) WATCH GTC ←	Same as left (except prob. no abort.)

"O.F."=Overflow,to many-

CONTINUING ←

OCCURRENCE OF?

01109 DELAY ROUT. O.V. | BALLOUT

01201 EXECT. O.F.(VAC) | "

01202 EXECT. O.F.(JOBS) | "

DUTY CYCLE MAY DEGRADE
PGNCS (AGS CONTROL MAY
HELP-SF E BELOW)
[WATCH FOR OTHER CUES]

PGNCS CONDITION UNKNOWN
DSKY MAY BE LOCKED UP,

DUTY CYCLE MAY DEGRADE
PGNCS (AGS CONTROL MAY
HELP-SF E BELOW)
[WATCH FOR OTHER CUES]

DSKY MAY BE LOCKED UP,

BAILOUT No available core sets

ALARM_AND_ABORT.AGC

BAILOUT INHINT

CA Q

TS ALMCADR

INDEX Q

CAF Ø

TC BORTENT

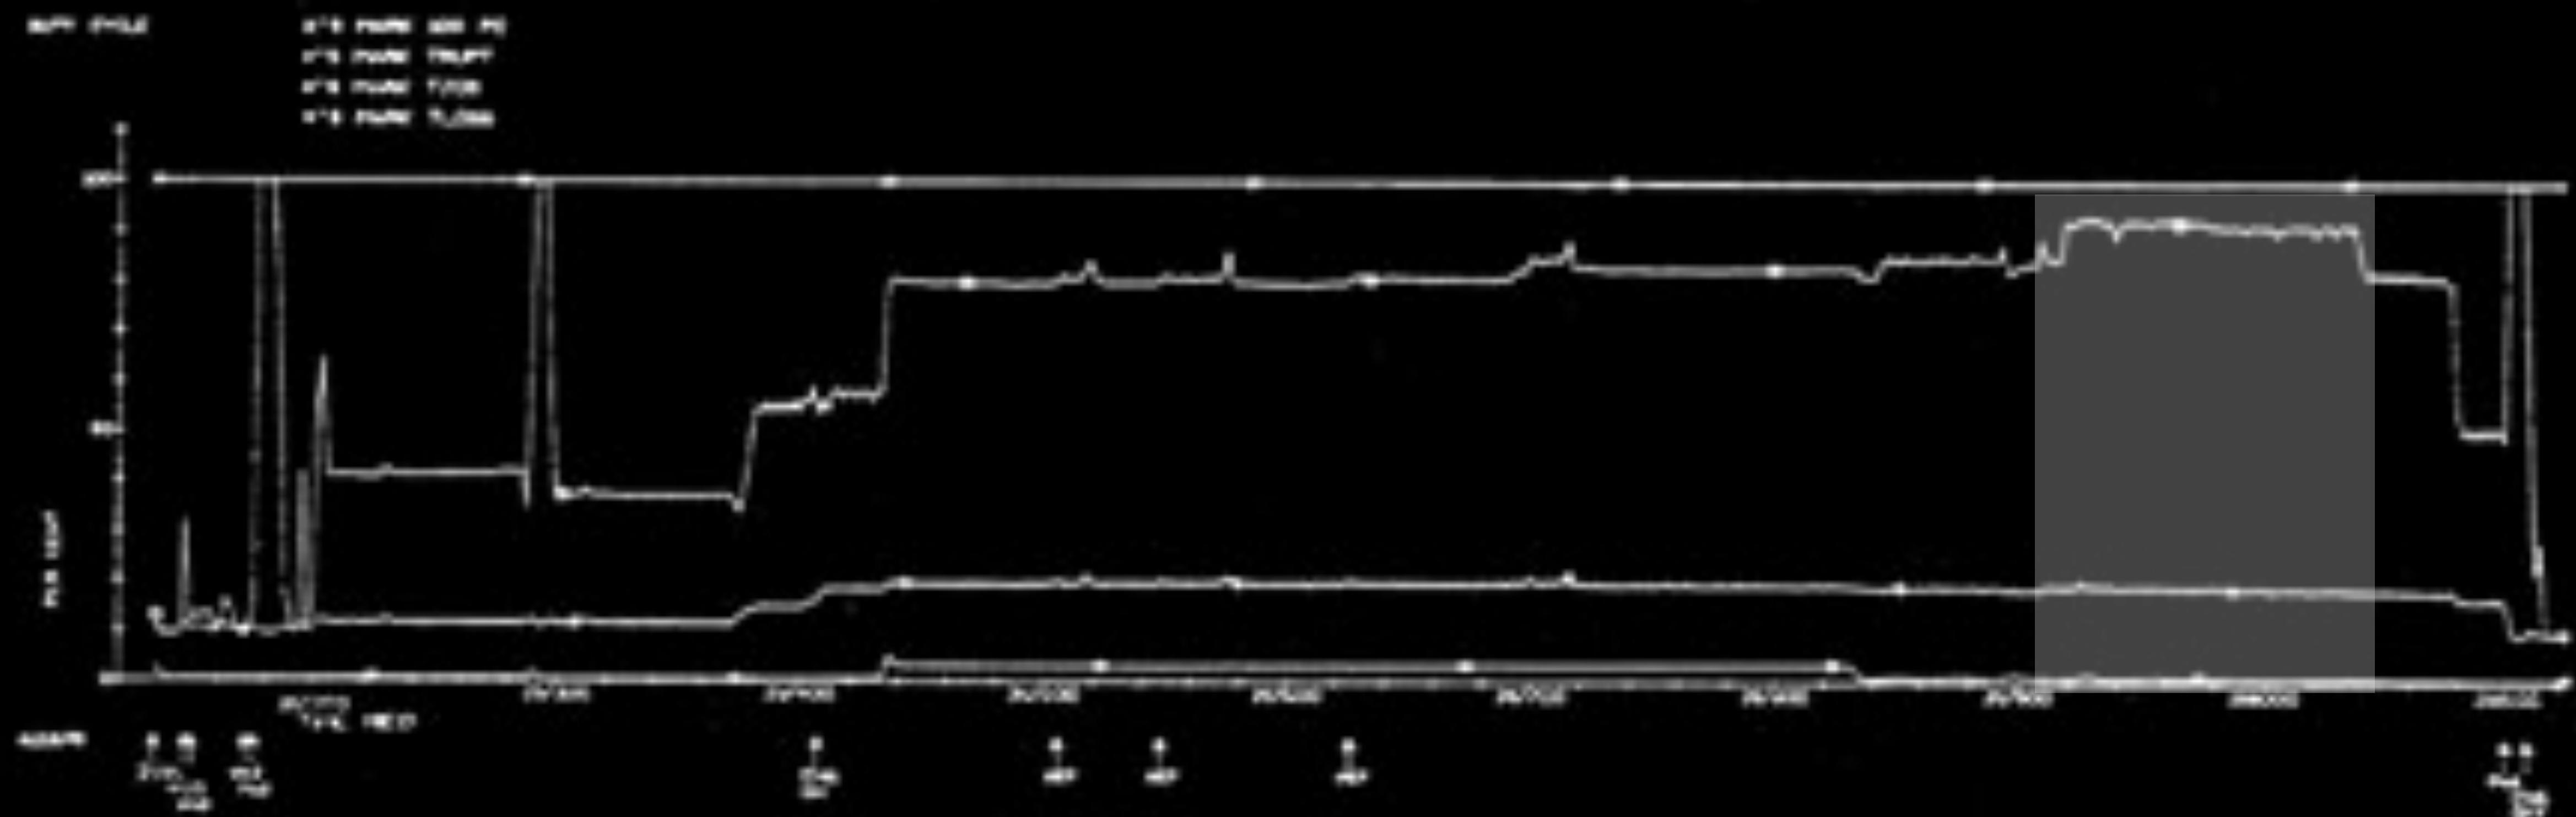
PERFORMANCE EVALUATION OF A MULTIPROCESSOR
IN A
REAL TIME ENVIRONMENT

by

Jaynarayan H. Lala

at the
MASSACHUSETTS INSTITUTE OF TECHNOLOGY
FEBRUARY 1976

The largest computational loads maintained on the Apollo Guidance Computer during the mission were imposed during the descent phase of the lunar landing. The programs were analyzed to determine their execution times in terms of the number of instructions executed, data requirements, and cyclic execution rates for periodic jobs.



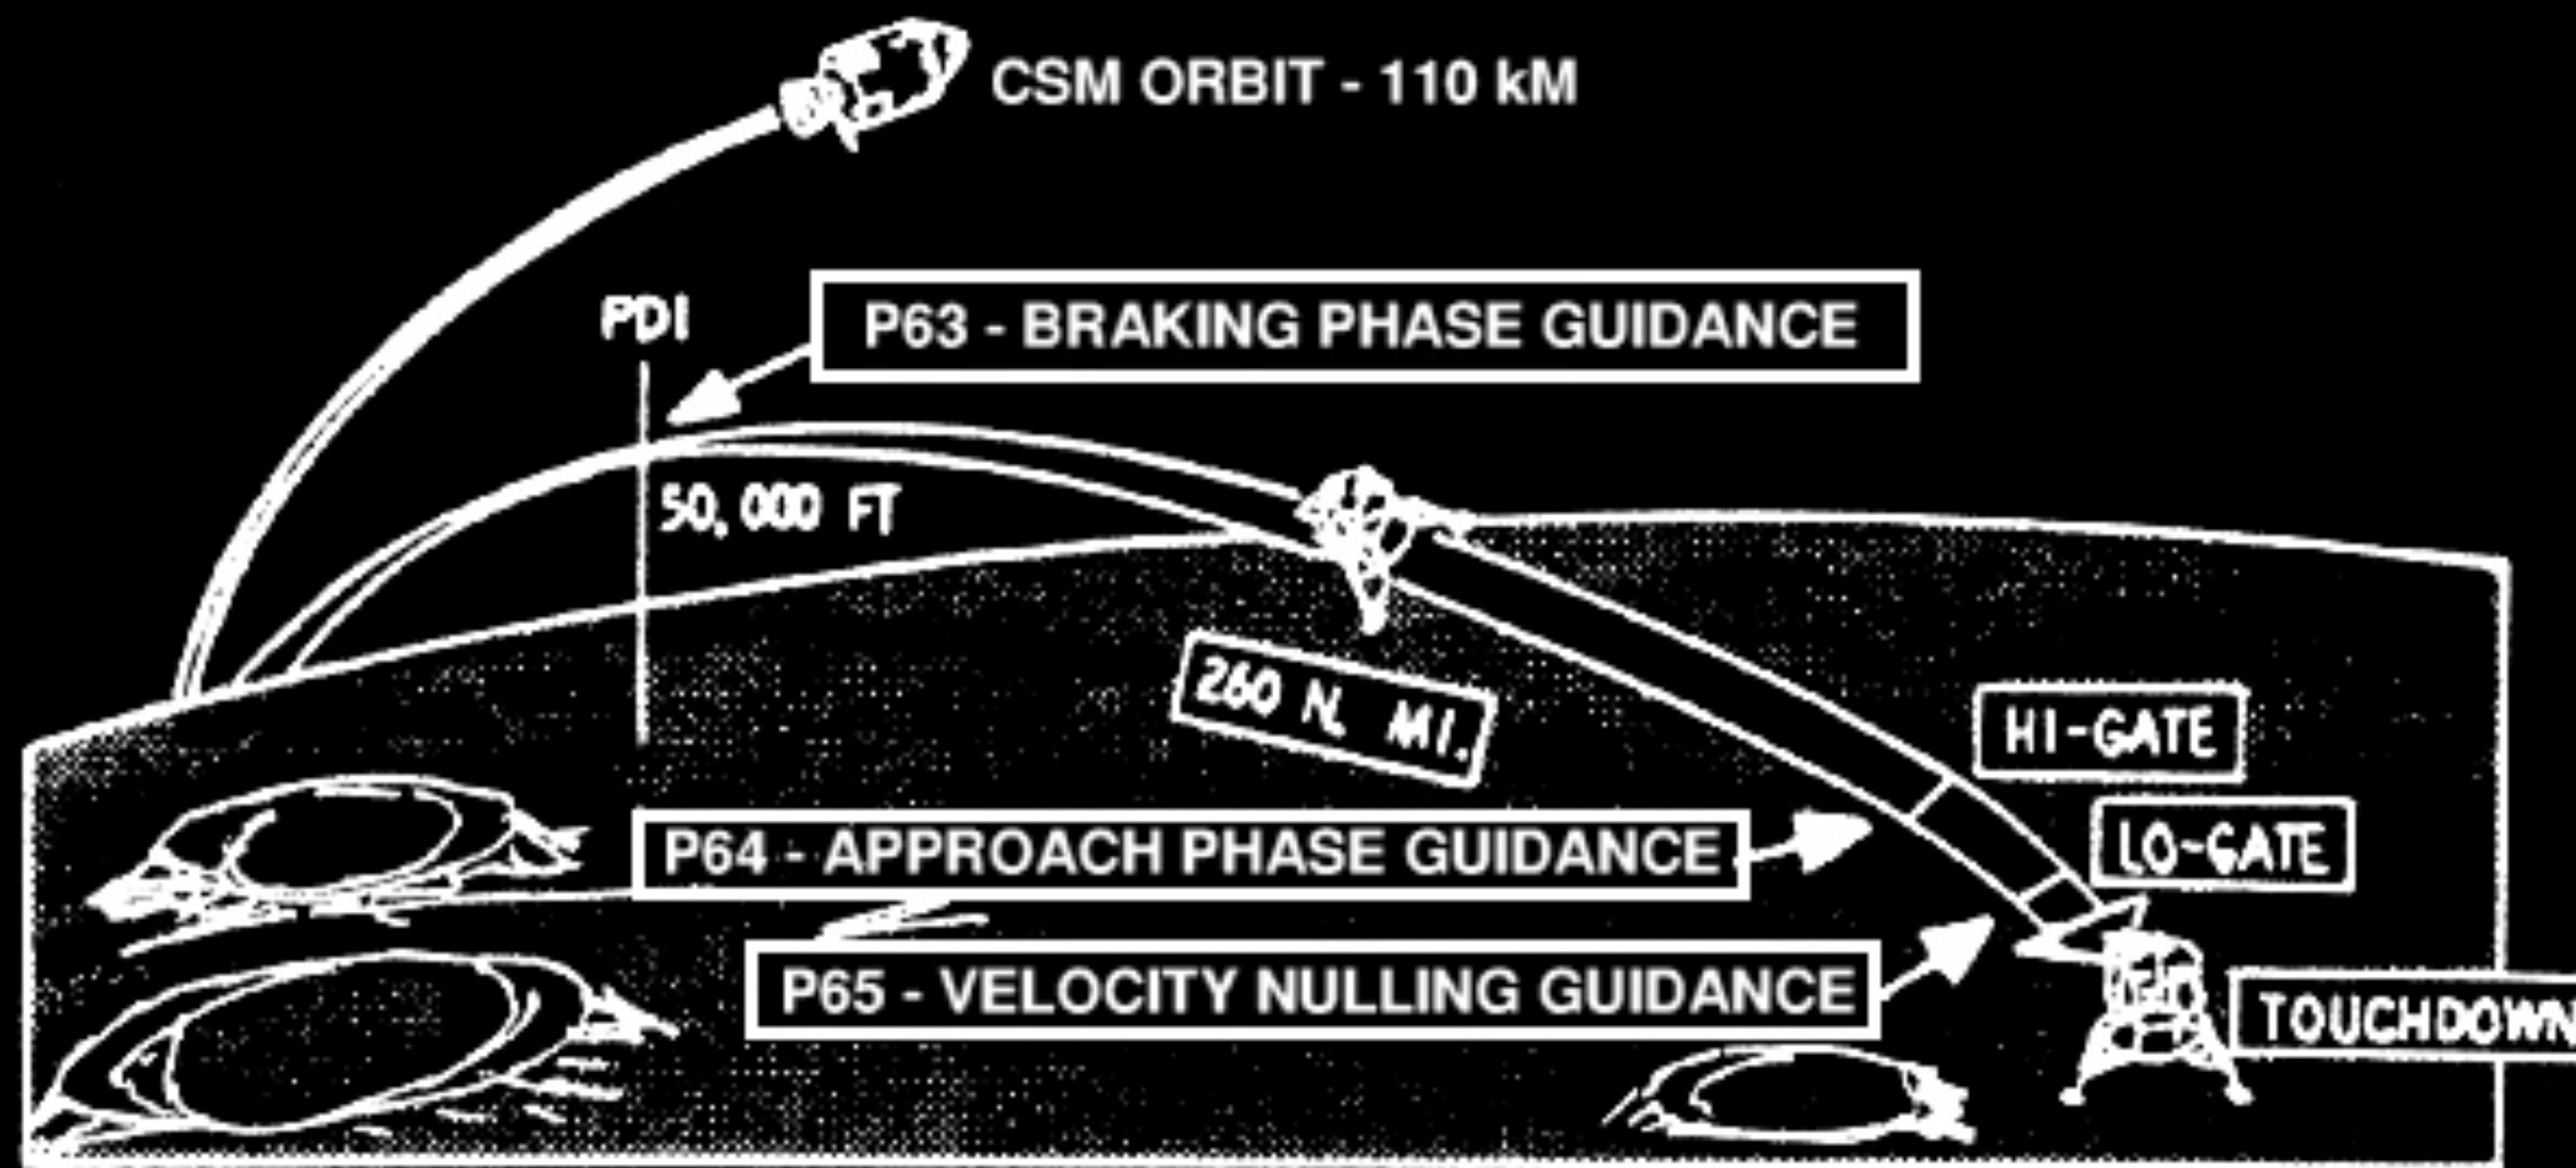


Fig. 3. Operational phases of powered descent. Design criteria: P 63 - Braking phase (PD1 to Hi-Gate), efficient reduction of orbital velocity. P 64 - Final approach phase (Hi-Gate to Lo-Gate), crew visibility (safety and site assessment). P 65 - Landing phase (Lo-Gate to touch-down), manual control takeover.

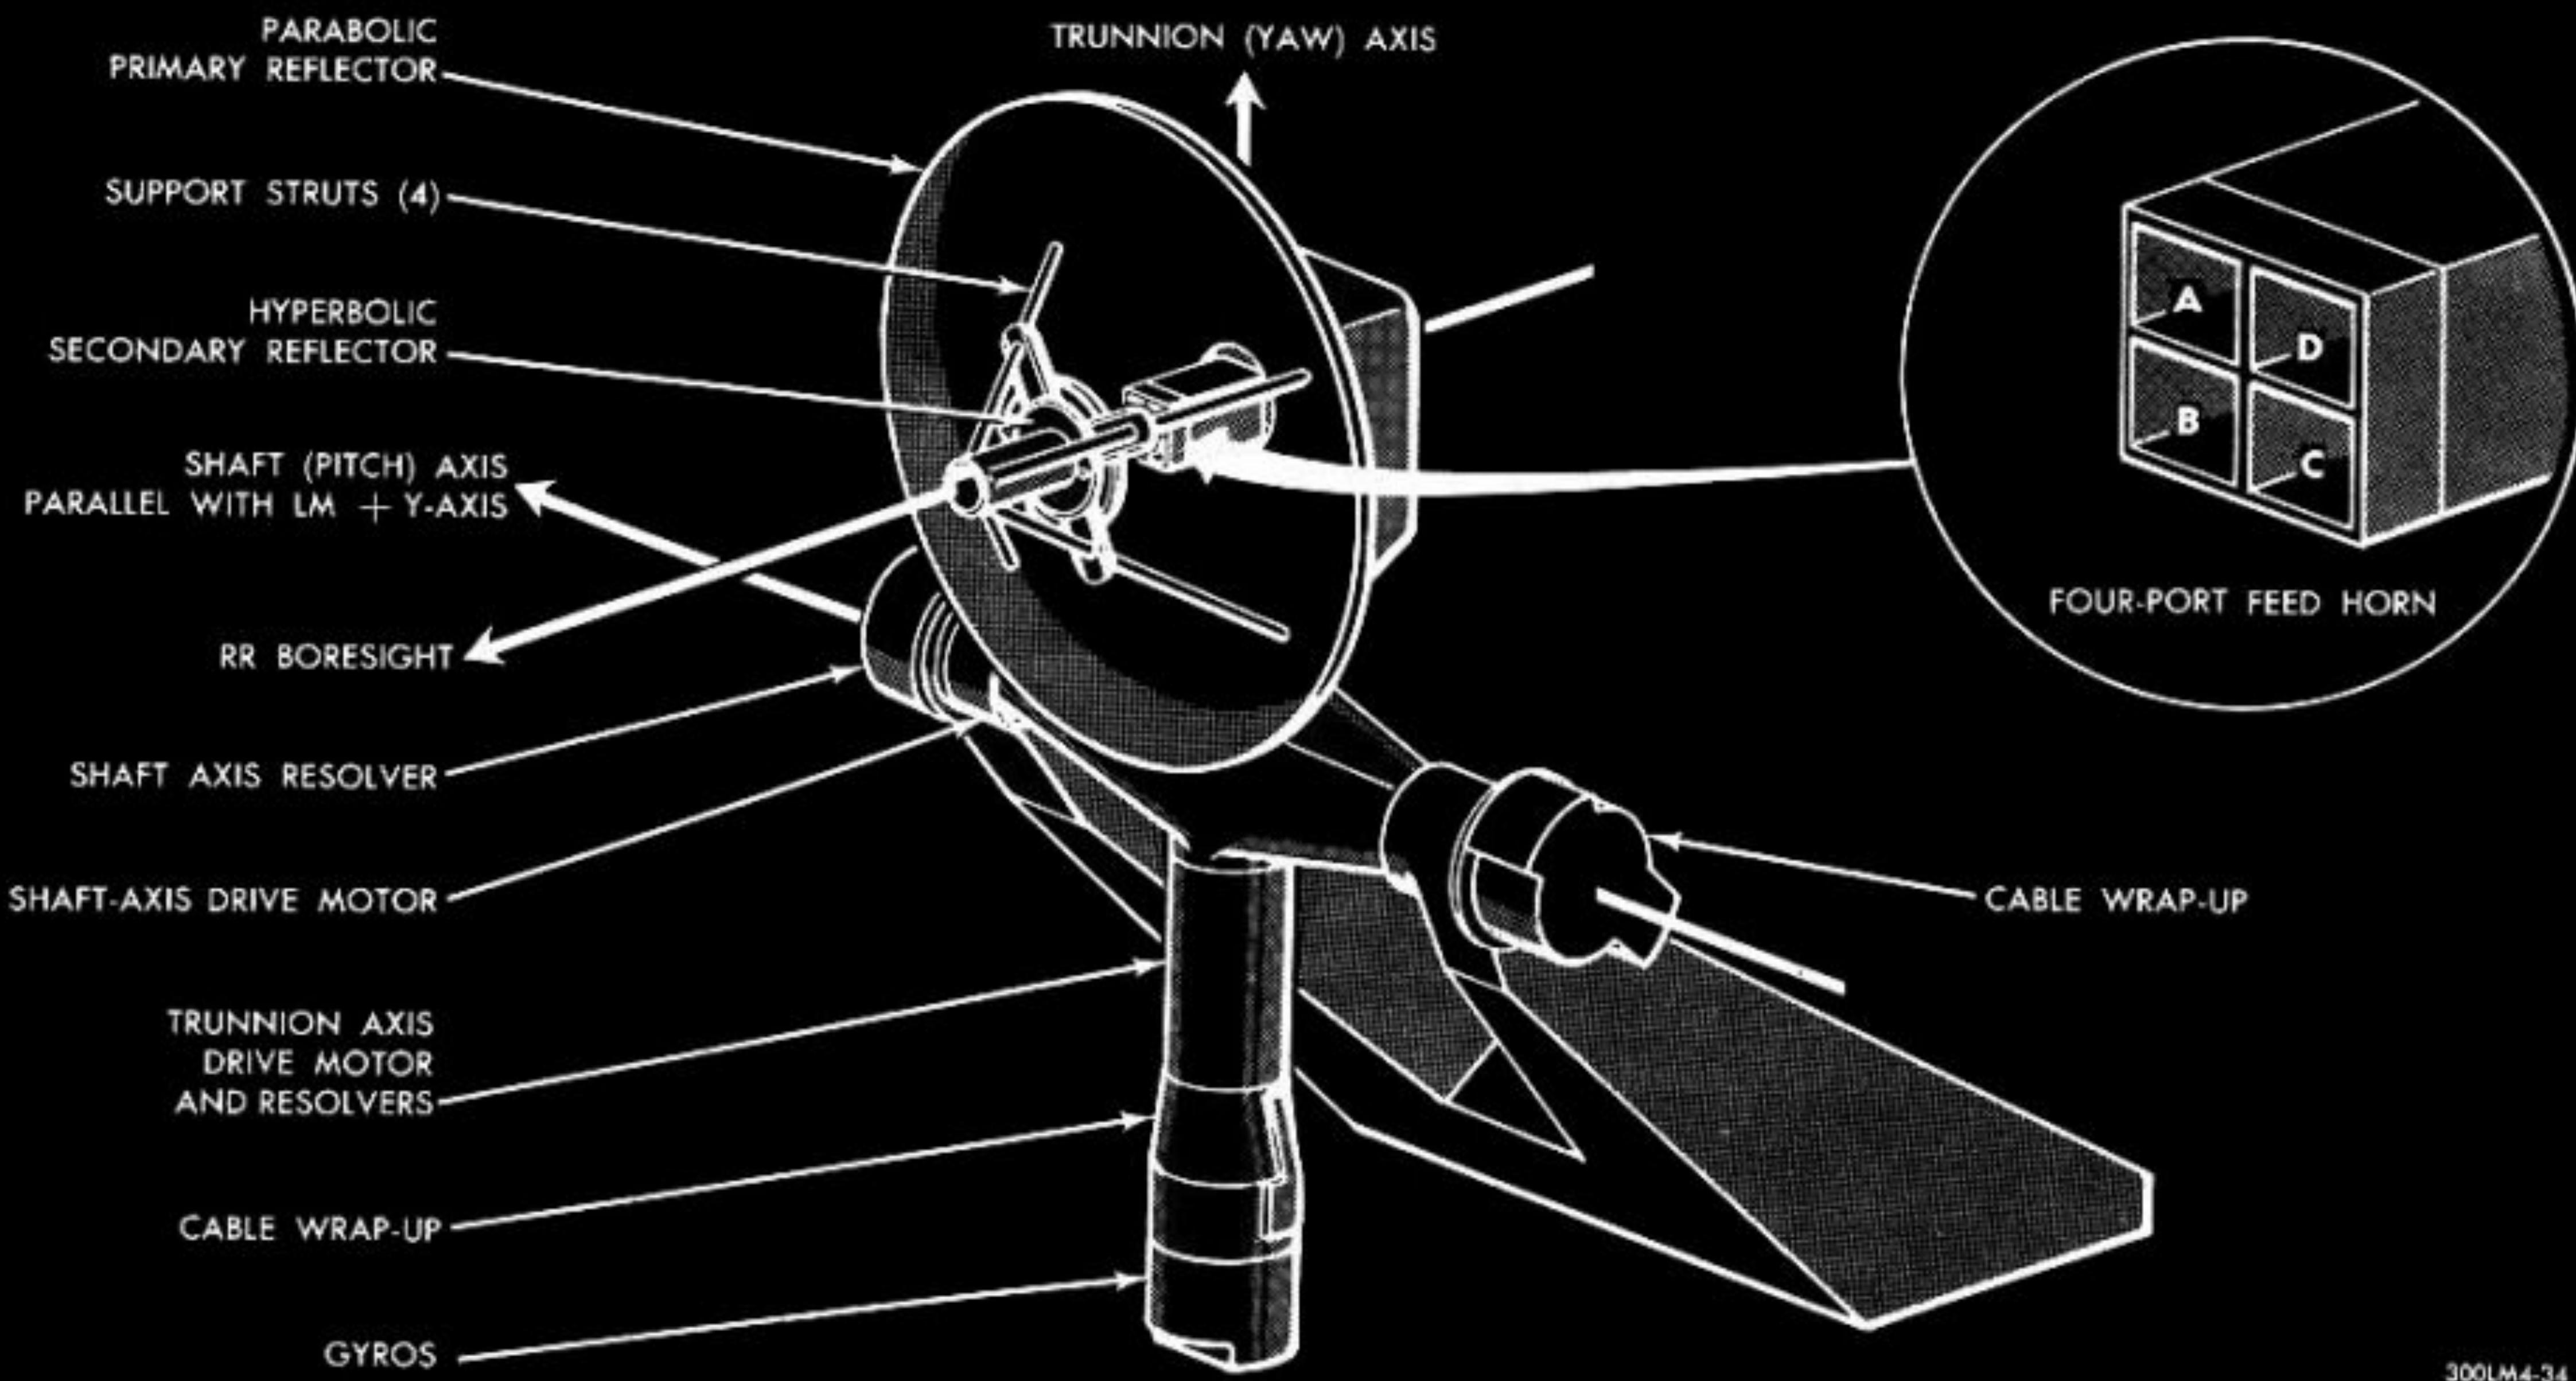


Figure 2.2-24. Rendezvous Radar Antenna Assembly

Apollo:

the ultimate

"TESTING IN PROD"

321 (239 unique) GCC-8 backend passes in execution order

```
all_lowering_passes
 1. warn_unused_result
 2. diagnose_omp_blocks
 3. diagnose_tm_blocks
 4. lower_omp
 5. lower_cf
 6. lower_tm
 7. refactor_sh
 8. lower_sh
 9. build_cfg
10. warn_function_return
11. expand_omp
12. sprintf_length[1]
13. valloca[1]
14. build_cgraph_edges

all_small_ipa_passes
15. ipa_free_lang_data
16. ipa_fn_and_var_visibility
17. ipa_chkp_vernioning
18. ipa_chkp_early_produce_thunks
build_ssa_passes
19. fixmp_cfg[1]
20. build_ssa
21. warn_nonsnull_compare
22. usan
23. early_warn_uninitialized
24. nothrow
25. rebuild_cgraph_edges[1]

chkp_instrumentation_passes
26. fixmp_cfg[2]
27. chkp
28. rebuild_cgraph_edges[2]

local_optimization_passes
29. fixmp_cfg[3]
30. rebuild_cgraph_edges[3]
31. local_fn_summary[1]
32. early_inline
all_early_optimizations
33. remove_cgraph callees edges[1]
34. object_sizes[1]
35. ccp[1]
36. forvprop[1]
37. early_thread_jumps
38. sra_early
39. build_aliases
40. fre[1]
41. early_vrp
42. merge_phi[1]
43. dce[1]
44. cd_dce[1]
45. early_ipa_sra
46. tail_recursion[1]
47. convert_switvh
48. cleanup_sh[1]
49. profile
50. local_pure_const[1]
51. split_functions
52. strip_predict_hints[1]

all_late_ipa_passes
53. release_ssa_names
54. rebuild_cgraph_edges[4]
55. local_fn_summary[2]
ipa_cacc
56. ipa_pta[1]
ipa_cacc_kernels
57. ch[1]
58. fre[2]
59. lim[1]
60. dominator[1]
61. dce[1]
62. parallelize_loops[1]
63. expand_omp_ssa[1]
64. rebuild_cgraph_edges[6]

all_regular_ipa_passes
65. target_clone
66. ipa_chkp_produce_thunks
67. ipa_auto_profile
68. ipa_tree_profile
69. feedback_split_functions
70. ipa_free_fn_summary[1]
71. ipa_increase_alignment
72. ipa_tm
73. ipa_lower_ssa
ipa_cacc
74. ipa_whole_program_visibility
75. ipa_profile
76. ipa_icf
77. ipa_devirt
78. ipa_cp
79. ipa_cctor_merge
80. ipa_haa
81. ipa_fn_summary
82. ipa_inline
83. ipa_pure_const
84. ipa_free_fn_summary[2]
85. ipa_reference
86. ipa_single_use
87. pass_ipa_consts
ipa_cacc
88. materialize_all_clones
89. ipa_pta[2]
90. cmp_ssa_clone
91. fixup_cfg[4]
92. lower_sh_dispatch
93. cacc_device_lower
94. cmp_device_lower
95. cmp_target_link
96. remove_cgraph callees edges[2]
97. strip_predict_hints[2]
98. ccp[2]
99. post_ipa_varn[1]
100. complete_umrolli

all_optimizations
101. backprop
102. phi_prop
103. forprop[2]
104. object_sizes[2]
105. build_aliases
106. return_slot
107. fre[3]
108. merge_phi[2]
109. thread_jumps[1]
110. vrp[1]
111. chkp_opt
112. dce[2]
113. stdarg
114. call_dce
115. cselim
116. copy_prop[1]
117. tree_ifcombine
118. merge_phi[3]
119. phi_opt[1]
120. tail_recursion[2]
121. ch[2]
122. lower_complex[1]
123. sra
124. thread_jumps[2]
125. dominator[2]
126. isolate_erroneous_paths
127. phi_only_cprop[1]
128. dce[2]
129. reassoc[1]
130. dce[3]
131. forprop[3]
132. phi_opt[2]
133. ccp[3]
134. cse_sincos
135. optimize_bwavp
136. laddress
137. lim[2]
138. valloca[2]
139. pre
140. sink_code
141. sancov[1]
142. asan[1]
143. tsan[1]
144. dce[4]
145. fix_loops
tree_loop
146. tree_loop_init
147. tree_unswitch
148. sancov_cprop
149. loop_split
150. loop_jam
151. cd_dce[2]
152. iv_canon
153. loop_distribution
154. limterchange
155. copy_prop[2]
graphite
156. graphite_transforms
157. lim[3]
158. copy_prop[3]
159. dce[5]

all_optimizations_g
160. parallelize_loops[2]
161. expand_omp_ssa[2]
162. ch_vect
163. if_conversion
164. vectorize
165. dce[6]
166. predcom
167. complete_unroll
168. slp_vectorize[1]
169. loop_prefetch
170. iv_optimize
171. lim
172. tree_loop_dom
tree_no_loop
173. slp_vectorize[2]
174. simduid_cleanup[1]
175. lower_vector_ssa[1]
176. case_reciprocals
177. sprintf_length[2]
178. reassoc[2]
179. strength_reduction
180. split_paths
181. tracer
182. thread_jumps[3]
183. dominator[3]
184. strien
185. thread_jumps[4]
186. vrp[2]
187. warn_restrict
188. phi_only_cprop[2]
189. dce[3]
190. cd_dce[3]
191. forprop[4]
192. phi_opt[3]
193. fold_builtin[1]
194. optimize_widening_mil
195. store_merging
196. tail_calls
197. dce[7]
198. split_crit_edges[1]
199. late_warn_uninitialized[1]
200. uncprop[1]
201. local_pure_const[2]
loop2
202. remove_cgraph callees edges[3]
203. strip_predict_hints[3]
204. lower_complex[2]
205. lower_vector_ssa[2]
206. ccp[4]
207. post_ipa_varn[2]
208. object_sizes[3]
209. fold_builtin[2]
210. sprintf_length[3]
211. copy_prop[4]
212. dce[5]
213. sancov[2]
214. asan[2]
215. tsan[2]
216. split_crit_edges[2]
217. late_warn_uninitialized[2]
218. uncprop[2]
219. local_pure_const[3]
tm_init
220. tm_mark
221. tm_nsopt
222. tm_sedges
223. simduid_cleanup[2]
224. vtable_verify
225. lower_vaang
226. lower_vector
227. lower_complex_00
228. sancov_00
229. lower_switch
230. asan_00
231. tsan_00
232. sancov
233. cleanup_sh[2]
234. lower_rexx
235. nrv
236. cleanup_cfg_post_optimizing
237. warn_function_noreturn
238. gen_hmail
239. expand
rest_of_compilation
240. instantiate_virtual_regs
241. into_cfg_layout_mode
242. jump[1]
243. lower_subreg[1]
244. df_initialize_opt
245. cse[1]
246. rtl_fvprop
247. rtl_cprop[1]
248. rtl_pre
249. rtl_hoist
250. rtl_cprop[2]
251. rtl_stores_motion
252. rtl_after_global_opts
253. rtl_ifcvt
254. reginfo_init
loop
255. rtl_loop_init
256. rtl_move_loop_invariants
257. rtl_unroll_loops
258. rtl_doloop
259. rtl_loop_dom
260. web
261. rtl_cprop[3]
262. cse[2]
263. rtl_dce[1]
264. rtl_fvprop_addr
265. inc_dce
266. initialize_regs
267. ud_rtl_dce
268. combine
269. if_after_combine
270. partition_blocks
271. outof_cfg_layout_mode
split_all_inmms
272. lower_subreg[2]
273. df_initialize_no_opt
275. stack_ptr_mod
276. mode_swapping
277. match_mm_constraints
278. sra
279. live_range_shrinkage
280. sched[1]
281. early_remat
282. ira
283. reload
postreload
284. postreload_cse
285. gcse
286. split_after_reload
287. rre
288. compare_slim_after_reload
289. branch_target_load_optimize[1]
290. thread_prologue_and_epilogue
291. rtl_dce[2]
292. stack_adjustments
293. jump[2]
294. duplicate_computed_gotos
295. sched_fusion
296. peephole
297. if_after_reload
298. regnames
299. cprop_hardreg
300. fast_rtl_dce
301. reorder_blocks
302. branch_target_load_optimize[2]
303. leaf_regs
304. split_before_sched2
305. sched[2]
stack_regs
306. split_before_regs
307. stack_regs_run
late_compilation
308. compute_alignments
309. variable_tracking
310. free_cfg
311. machine_reorg
312. cleanup_barriers
313. delay_slots
314. split_for_shorten_branches
315. convert_to_sh_region_ranges
316. shorten_branches
317. set_nothrow_function_flags
318. dwarf2_frame
319. final
320. df_finish
321. clean_state
```

**WE
are not
COMPILERS**

**ALL PERFORMANCE
WORK WITHOUT
MEASUREMENT is
PREMATURE**

MEASURE TWICE PR ONCE



**A BENCHMARK
is a SYNTHETIC MEASUREMENT
of the RESOURCES or TIME consumed
by a defined PROCEDURE**

A PROFILE
is a STEP BY STEP ACCOUNTING
of the RELATIVE CONSUMPTION
of RESOURCES by a procedure's
many SUBROUTINES



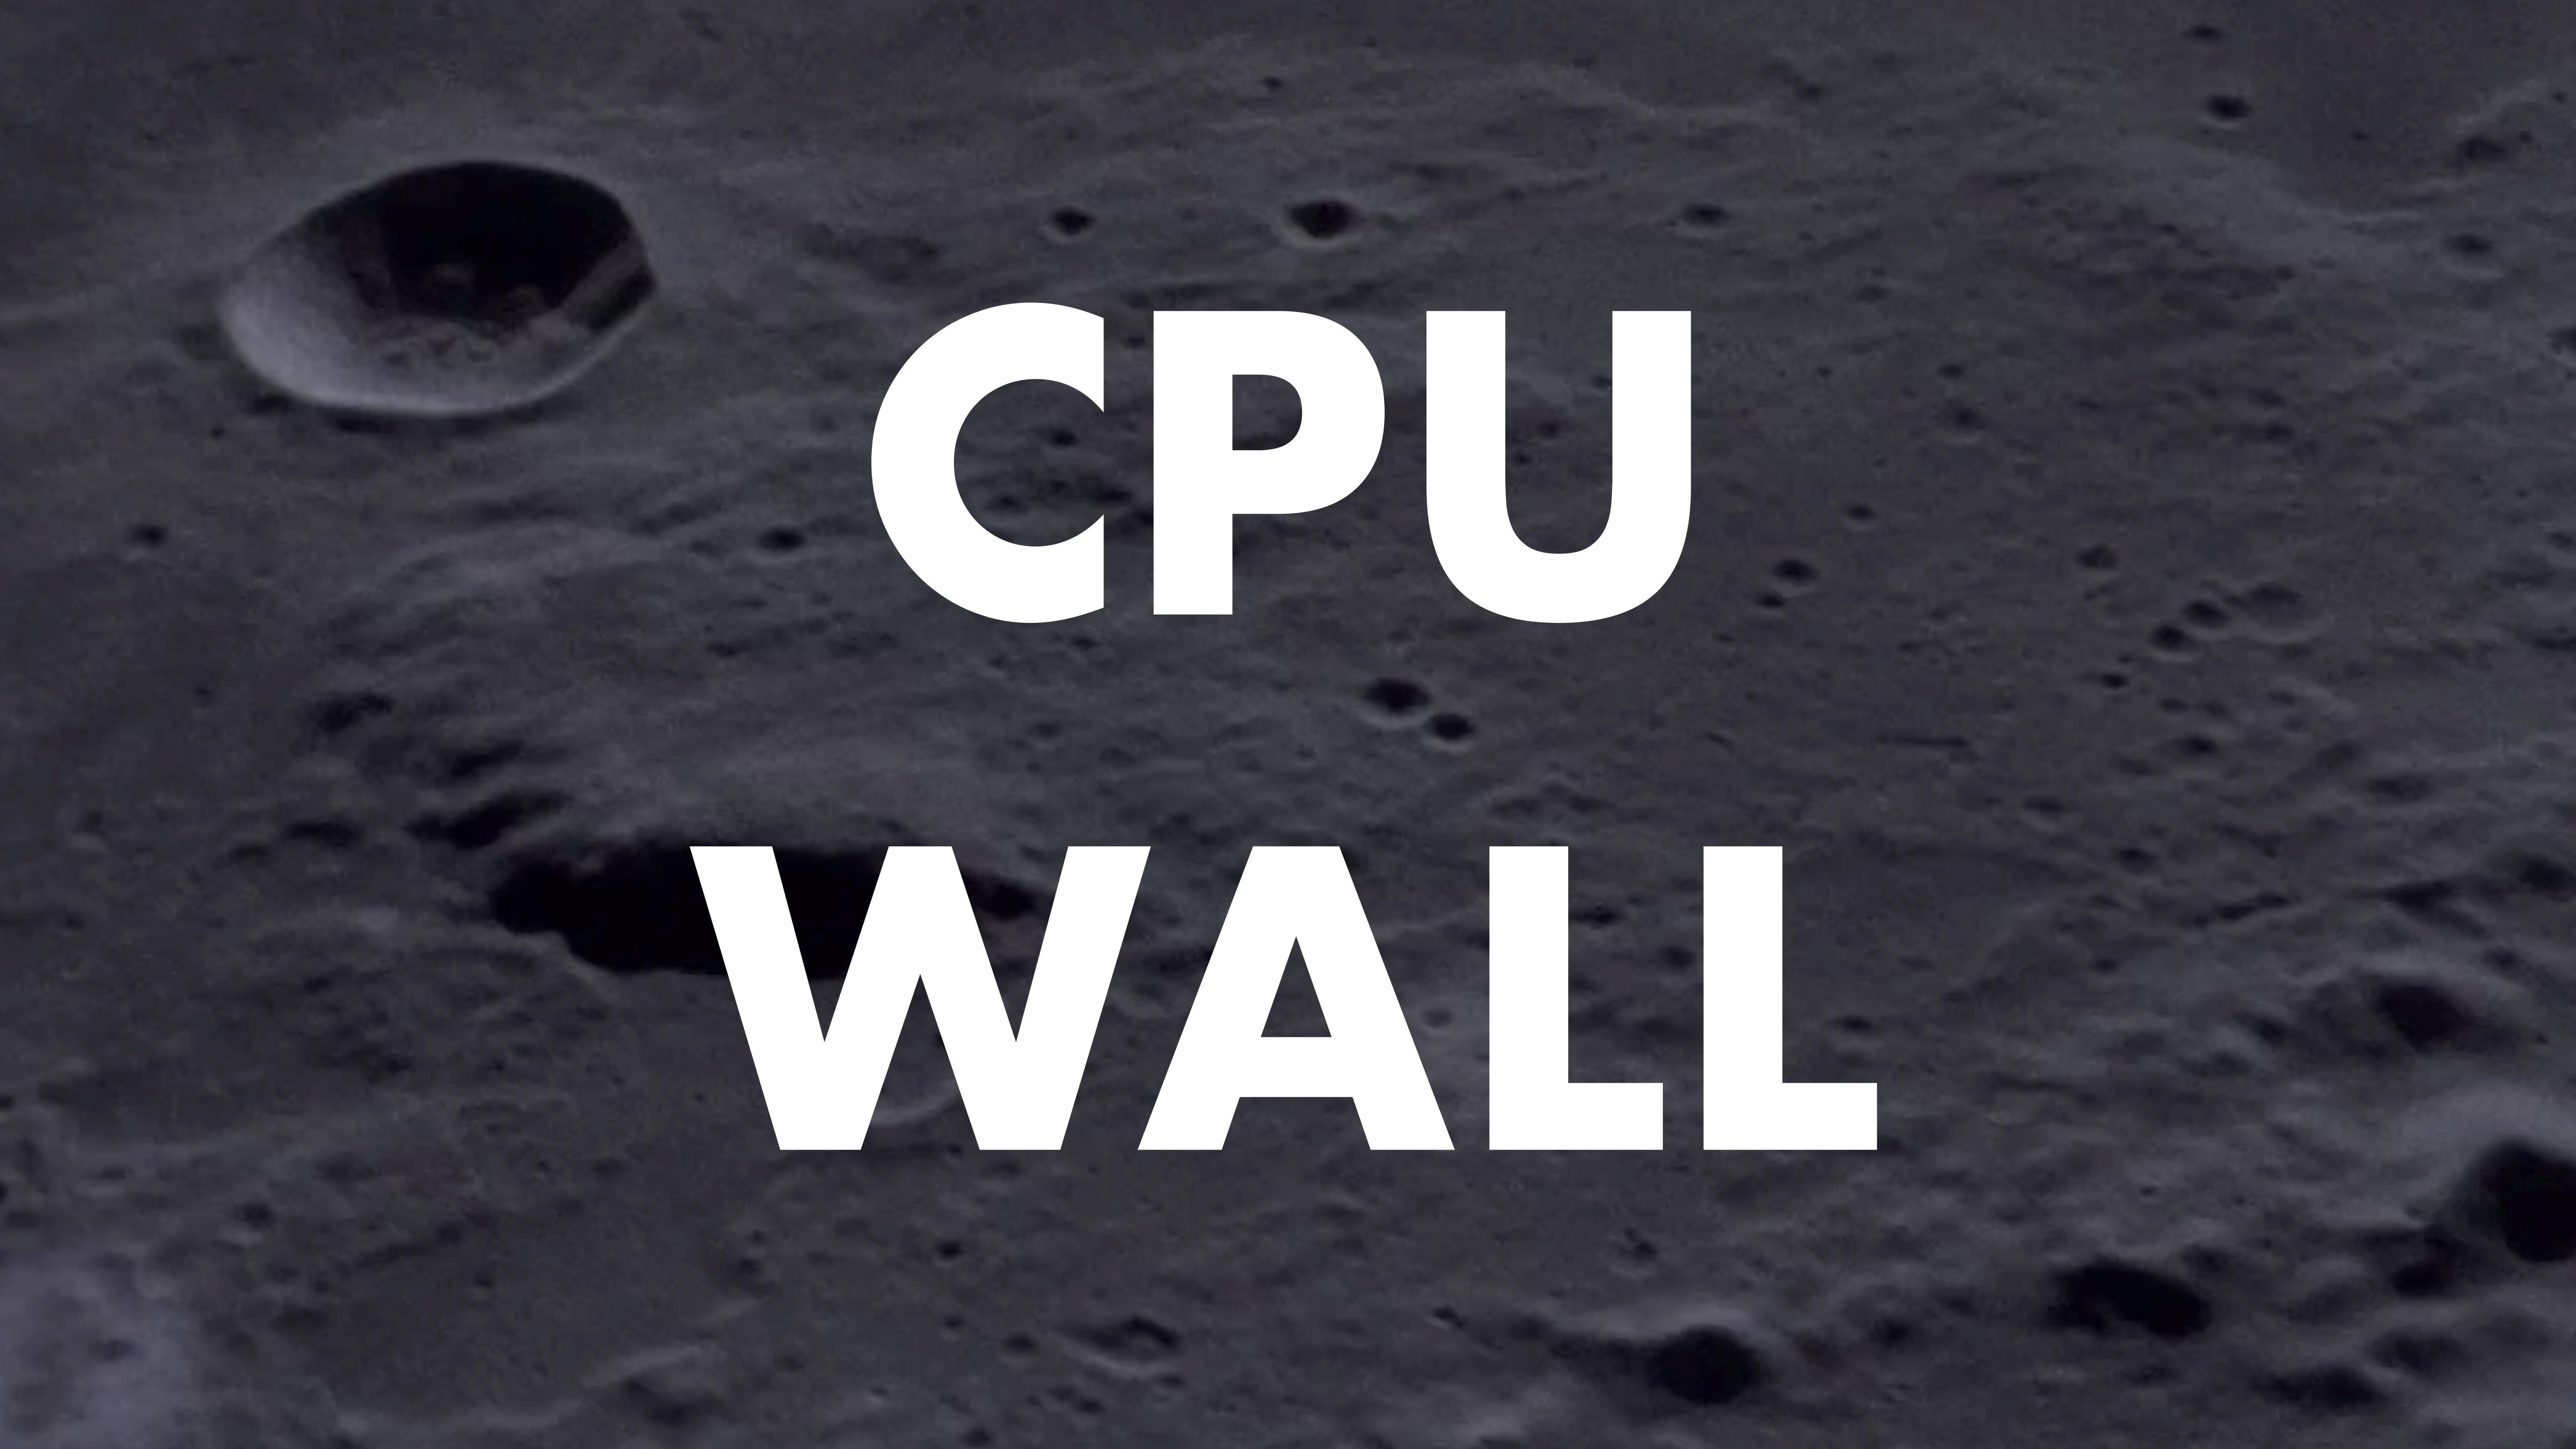


**READ PRODUCTION METRICS
PROFILE TO FIND HOTSPOTS
CREATE BENCHMARK
ITERATE
DEPLOY**

PROFILING
answers
“WHAT’S SLOW?”

PROFILING ENVIRONMENT

```
config.cache_classes = true
config.eager_load = true
config.public_file_server.enabled = true
config.assets.compile = false
config.assets.digest = true
config.active_record.migration_error = false
```



**CPU
WALL.**

```
def sleeper  
    sleep(4)  
end
```

```

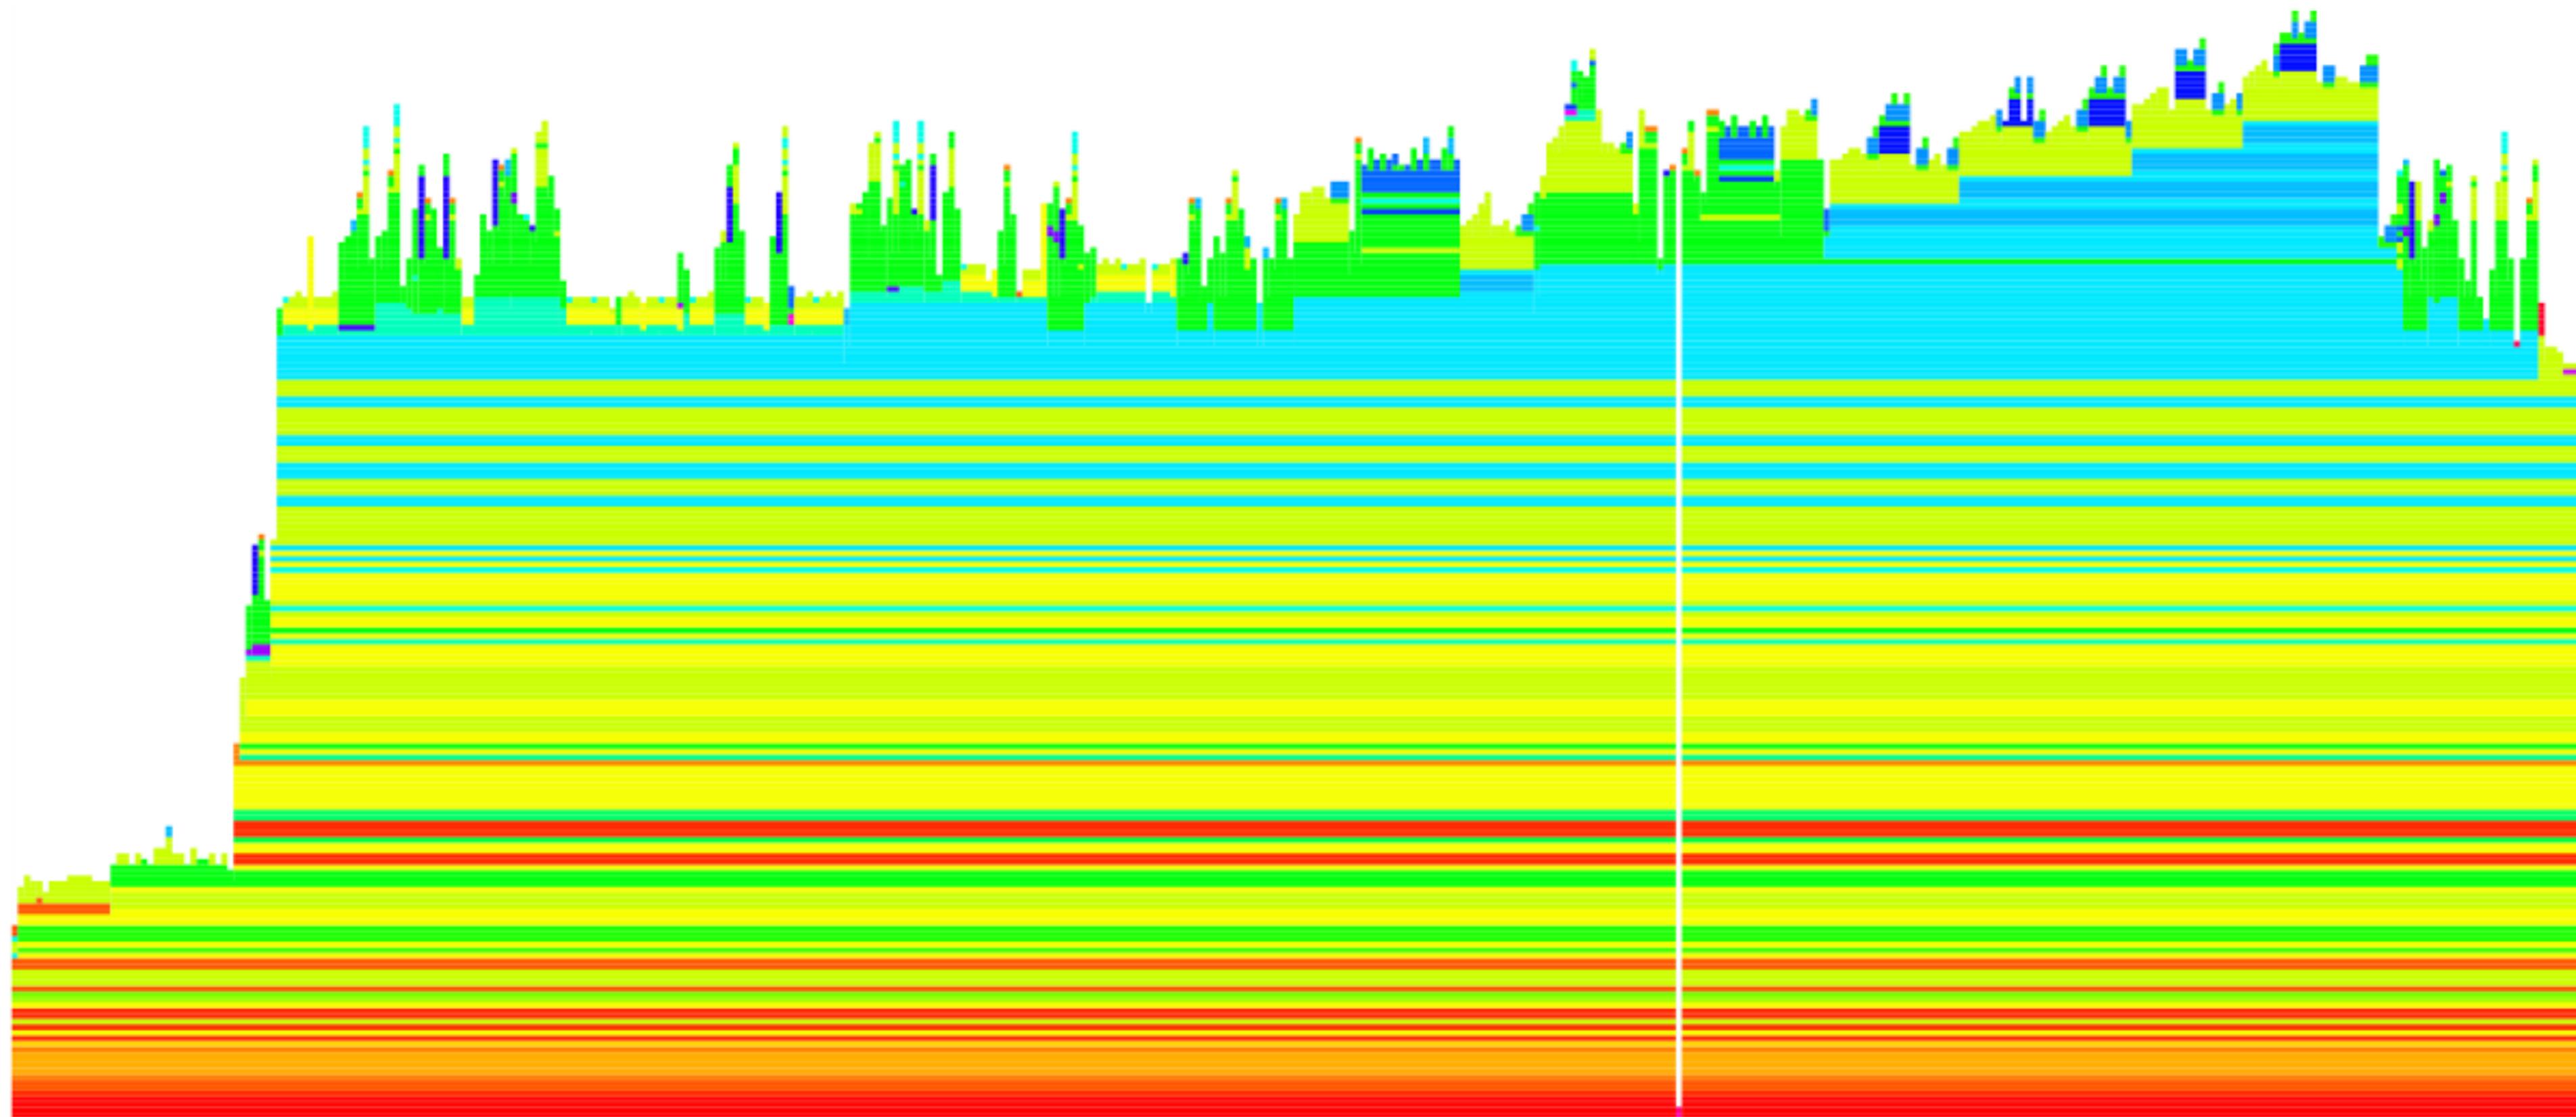
def sleeper
  sleep(4)
end

```

Measure	Mode: wall_time					
Thread	ID: 70311995300360					
Fiber I	D: 70311999772520					
Total:	4.00332					
Sort by	: self_time					
%self	total	self	wait	child	calls	name
100	4.003	4.003	0	0	1	Kernel#sleep
0	4.003	0.000	0	4.003	1	Global#[No method]
0	4.003	0.000	0	4.003	1	Object#sleeper

STATISTICAL TRACING

STACKPROF
RUBY-PROF



puma-2.12.2 (420 samples - 99.53%)	rack-1.6.4 (420 samples - 99.53%)	railties-4.2.6 (420 samples - 99.53%)	rack-mini-profiler-0.10.1 (420 samples - 99.53%)	flamegraph-0.9.5 (420 samples - 99.53%)	sentry-raven-0.14.0 (420 samples - 99.53%)	actionpack-4.2.6 (420 samples - 99.53%)	activesupport-4.2.6 (420 samples - 99.53%)	request_store-1.3.1 (419 samples - 99.29%)	quiet_assets-1.1.0 (419 samples - 99.29%)
web-console-2.2.1 (418 samples - 99.05%)	better_errors-2.1.1 (418 samples - 99.05%)	activerecord-4.2.6 (403 samples - 95.50%)	clearance-1.10.1 (383 samples - 90.76%)	jquery-fileupload-rails-0.4.7 (382 samples - 90.52%)	actionview-4.2.6 (381 samples - 90.28%)	eventstream (380 samples - 90.05%)	2.3.1 (377 samples - 89.34%)	active_model_serializers-c029c00758d5 (376 samples - 89.10%)	thread_safe-0.3.5 (112 samples - 26.54%)
binding_of_caller-0.7.2 (47 samples - 11.14%)	activemodel-4.2.6 (31 samples - 7.35%)	nilify_blanks-1.2.1 (27 samples - 6.40%)	did_you_mean-1.0.0 (27 samples - 6.40%)	arel-6.0.3 (13 samples - 3.08%)	acts_as_tenant-0.4.0 (8 samples - 1.90%)	paranoia-2.1.3 (6 samples - 1.42%)	friendly_id-5.1.0 (5 samples - 1.18%)	json-1.8.6 (4 samples - 0.95%)	kaminari-0.16.3 (2 samples - 0.47%)
pubnub-3.8.0 (1 sample - 0.24%)	event_machine.rb (1 sample - 0.24%)	eventmachine-1.2.0.1 (1 sample - 0.24%)	jsonapi-0.1.1.beta2 (1 sample - 0.24%)	tzinfo-1.2.2 (1 sample - 0.24%)					



Total Time (Time of Parent)

50% (25%)

98% (99.9%)

90% (99.9%)

**TRY DIFFERENT
VIEWS/FORMATS**

memory_profiler

rack-mini-profiler



**MICRO
MACRO
APP**

```
require 'benchmark/ips'

Benchmark.ips do |x|
  x.config(:time => 5, :warmup => 2)

  x.report("simple") { Raven.capture_exception(DIVIDE_BY_ZERO) }
  x.report("rails") { Raven.capture_exception(RAILS_EXC) }

  x.compare!
end
```

Warming up -----

simple	17.000	i/100ms
rails	6.000	i/100ms
lots o logs	15.000	i/100ms

Calculating -----

simple	156.087	(±26.3%)	i/s	-	714.000
rails	56.316	(±26.6%)	i/s	-	264.000
lots o logs	173.704	(±11.5%)	i/s	-	855.000

Comparison:

lots o logs:	173.7 i/s
simple:	156.1 i/s - same-ish: difference falls within error
rails:	56.3 i/s - 3.08x slower

Allocations -----					
	simple	863/32	alloc/ret	50/14	strings/ret
	rails	5895/14	alloc/ret	50/7	strings/ret
lots o logs	4853/410	alloc/ret	50/50	strings/ret	
Warming up -----					
	simple	17.000	i/100ms		
	rails	6.000	i/100ms		
lots o logs	15.000	i/100ms			
Calculating -----					
	simple	156.087	(±26.3%) i/s -	714.000	
	rails	56.316	(±26.6%) i/s -	264.000	
lots o logs	173.704	(±11.5%) i/s -	855.000		
Comparison:					
lots o logs:	173.7	i/s			
simple:	156.1	i/s - same-ish: difference falls within error			
rails:	56.3	i/s - 3.08x slower			

WARMUP
long enough
for STABLE RESULTS

BENCHMARK TIME
should be long enough
FOR LOW VARIANCE

BEWARE
the
LOOP

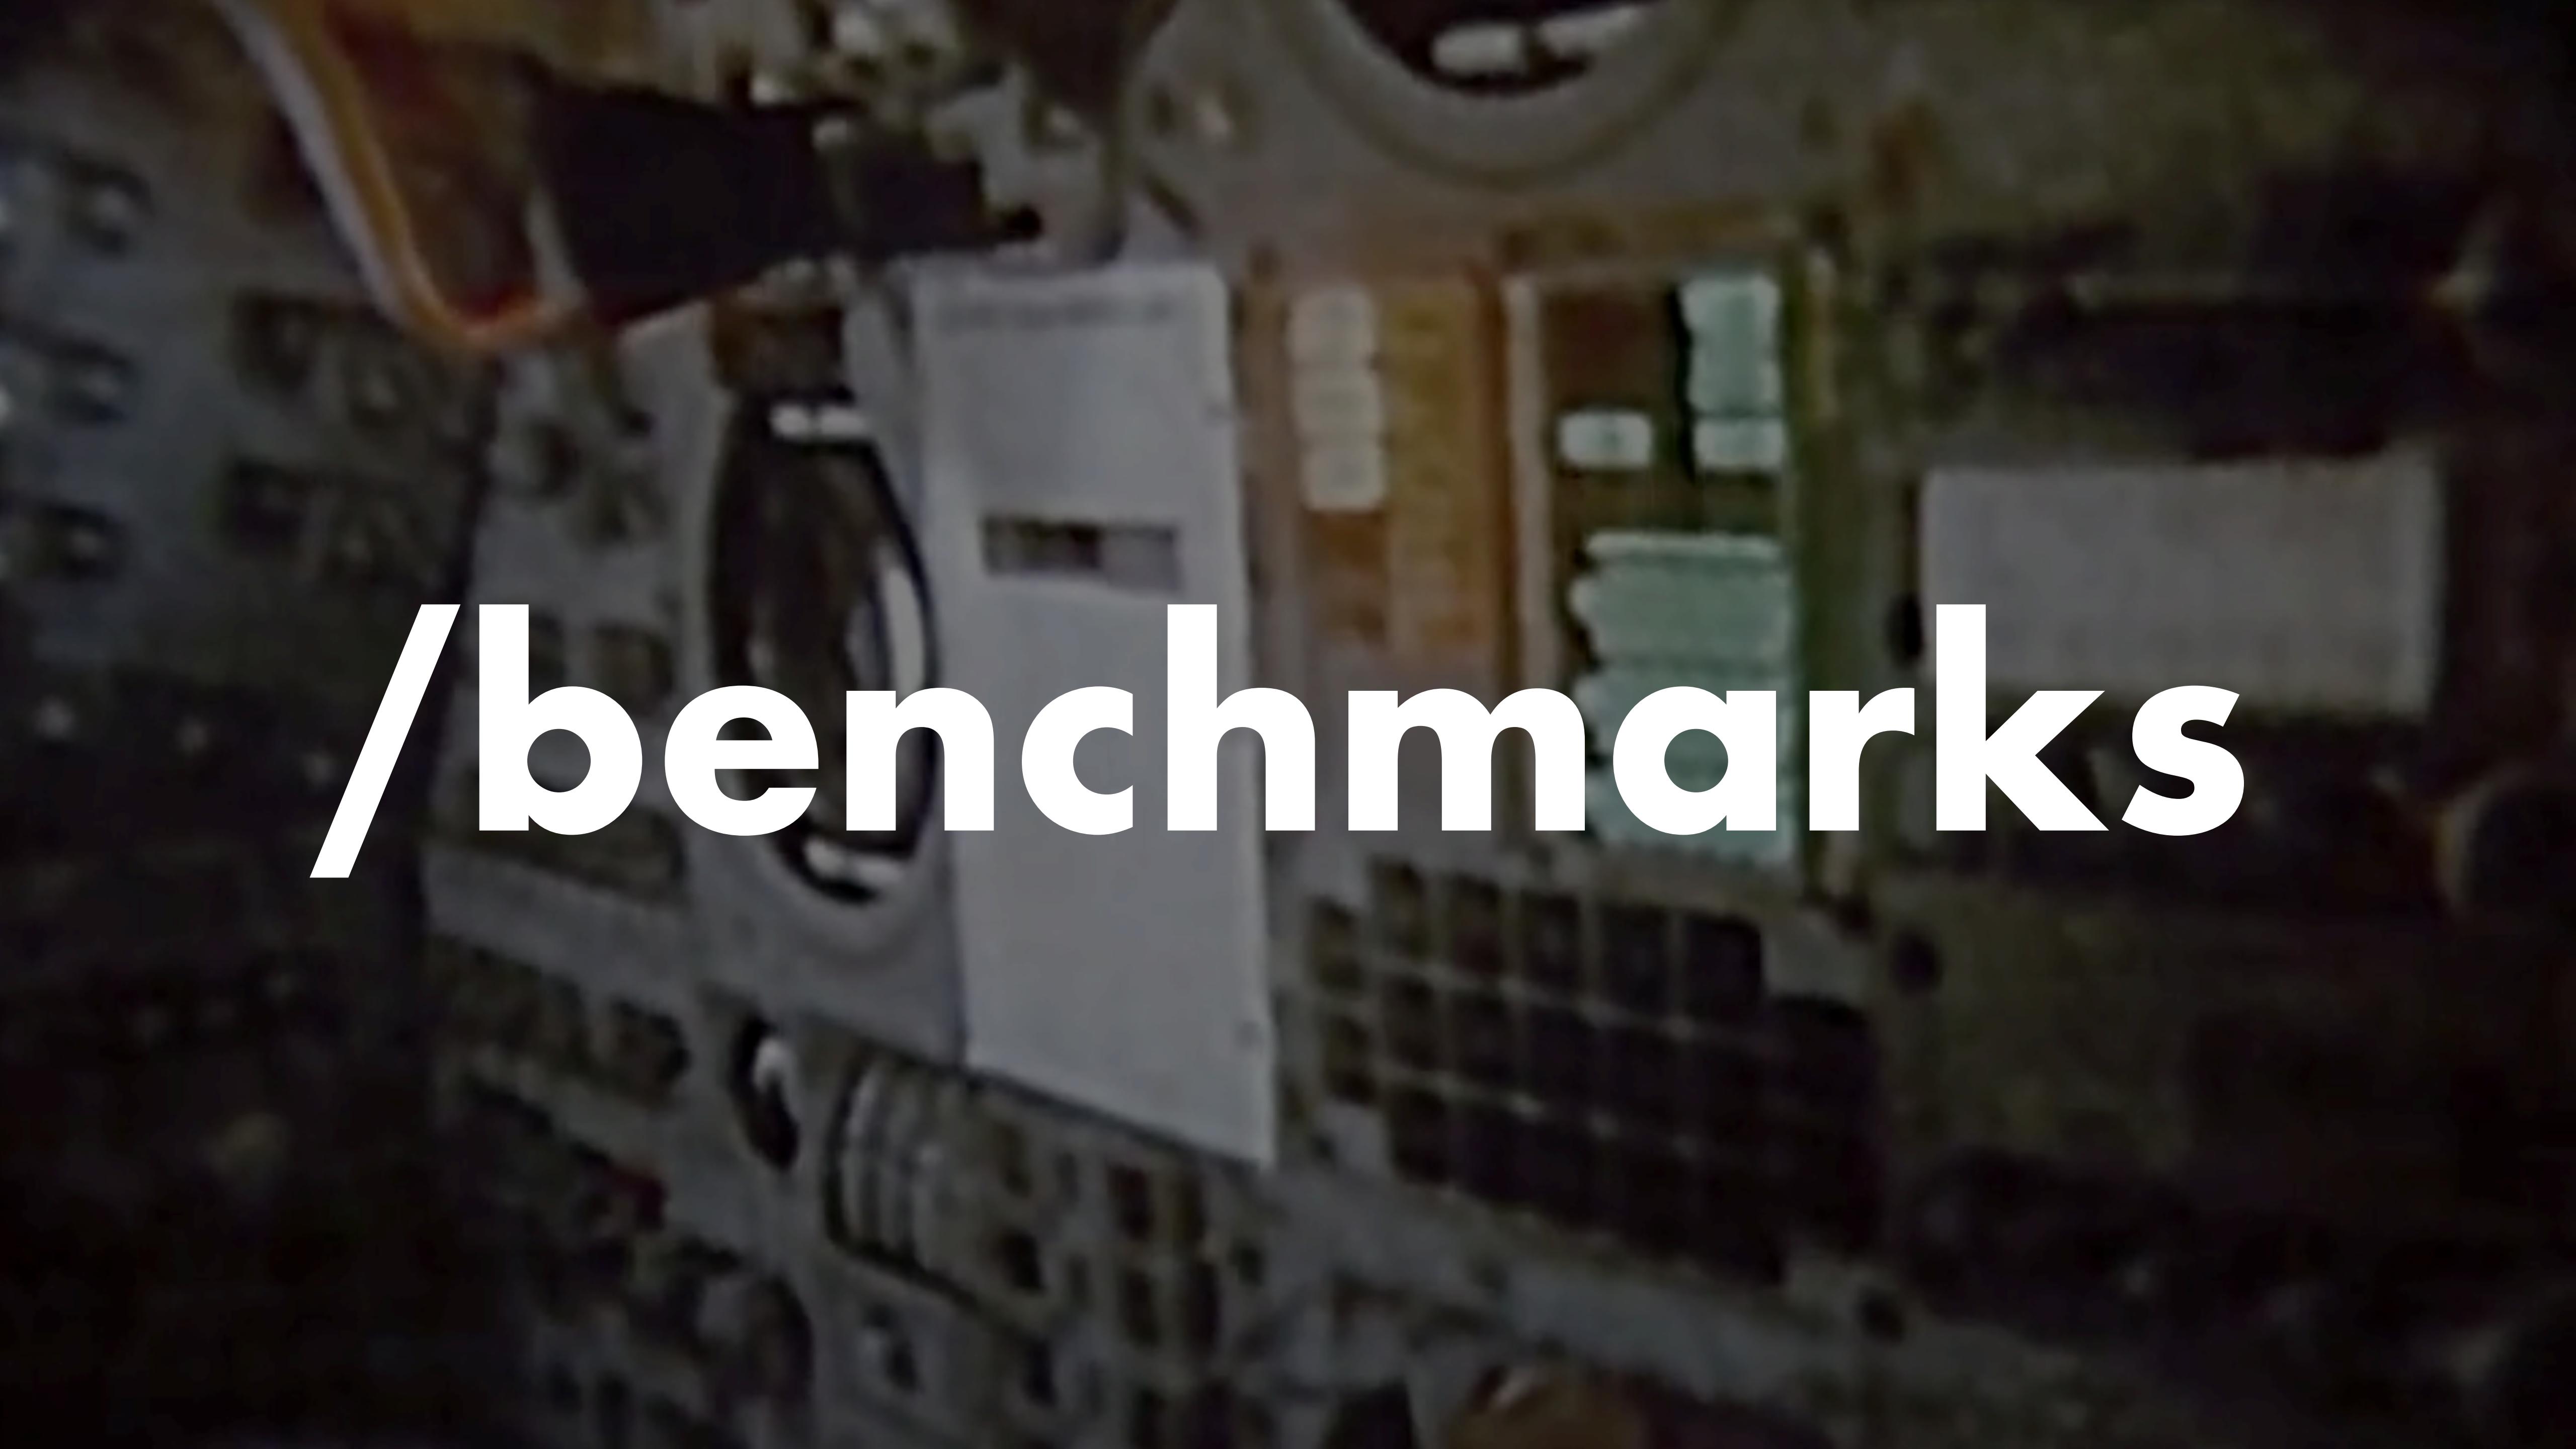
```
require 'benchmark/ips'
```

```
user = User.find(1)
```

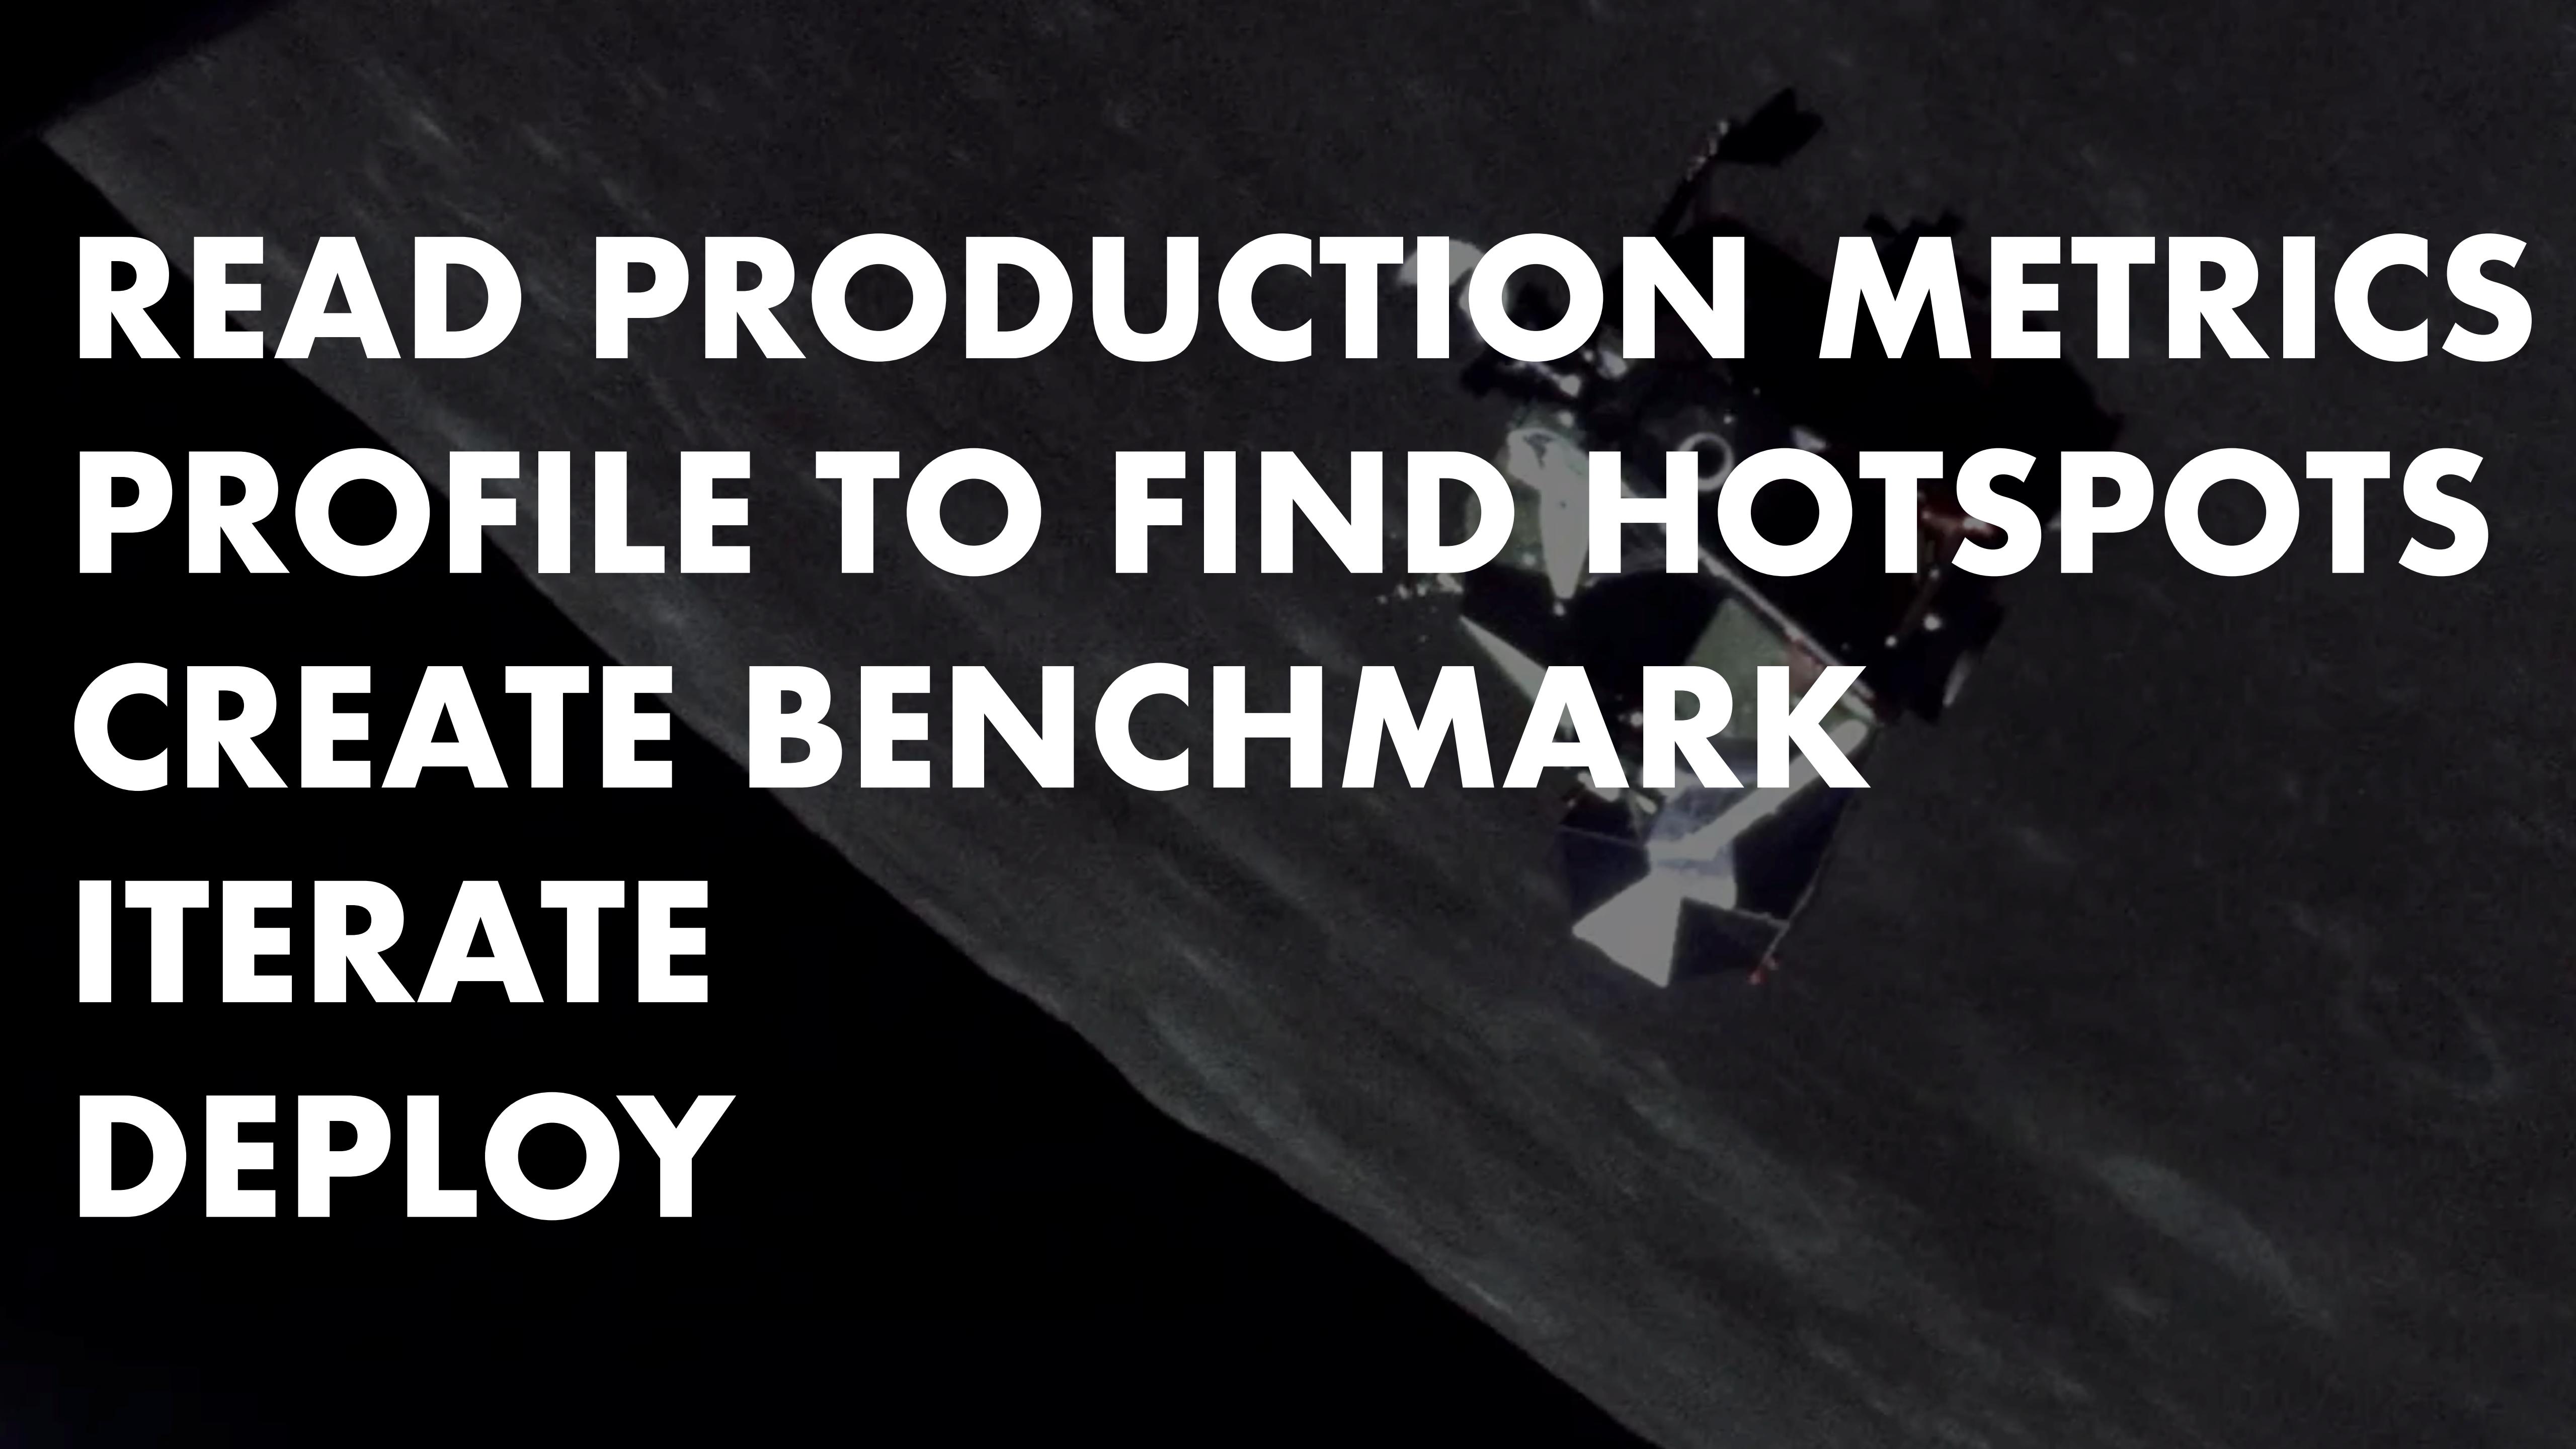
```
Benchmark.ips do |x|
  x.report("ar") { User.find(1).do_thing }
  x.report("ar2") { user.do_thing }
```

```
  x.compare!
```

```
end
```

A dark, grainy photograph of a city street at night. The scene is mostly black and grey, with some blurred lights from passing vehicles and a building visible in the background.

/benchmarks



**READ PRODUCTION METRICS
PROFILE TO FIND HOTSPOTS
CREATE BENCHMARK
ITERATE
DEPLOY**

FLIGHT PLAN

Why Measure?

Benchmarking vs. Profiling

Profiling Jumpstart

Benchmarking Jumpstart

**FAIL GRACEFULLY
DON'T PREDICT THE FUTURE
STUDY AND REPLICATE THE PAST**

speedshop

BENCHMARKING and PROFILING in RUBY

Railsconf 2019

